Machine Learning – Assignment 2

Multiclass classification of dry beans with five classifiers and a Streamlit evaluation app · M.Tech Artificial Intelligence & Machine Learning (AIML)

Name	**[Your full name]**
Student ID / Roll no.	**[Your BITS ID]**
Email	**[Your email]**
Programme / Course	M.Tech AIML · Machine Learning
Submission date	**[dd-mm-yyyy]**

1. GitHub repository link

<YOUR_GITHUB_REPO_URL>

2. Live Streamlit application link

<YOUR_STREAMLIT_APP_URL>

3. BITS Virtual Lab / Prayogshala screenshot

[INSERT YOUR BITS VIRTUAL LAB SCREENSHOT HERE]

It must visibly show: the Virtual Lab environment, the live app URL, the uploaded test CSV, the application heading and the results (see screenshots/README.md). Re-run this script with --screenshot to embed it.

4. README / report content

Dry Bean Classification — Machine Learning Assignment 2

M.Tech Artificial Intelligence & Machine Learning (AIML)

a. Problem Statement

Dry beans of different registered varieties look alike, and manual sorting is slow and subjective. The task is to build, evaluate and compare **five supervised classifiers** that predict the **bean variety** (`Class`, 7 categories) from 16 geometric measurements obtained from grain images. All models use the same data and the same stratified 80:20 train/test split and are scored with Accuracy, AUC, Precision, Recall, F1 and MCC. The best model is then exposed through an interactive Streamlit application.

b. Dataset Description

- **Name:** Dry Bean Dataset (UCI ID 602) — multiclass classification
- **Instances:** 13,611 (raw) → **13,543** after removing 68 exact duplicate rows
- **Input features:** 16 numeric (12 dimension features + 4 shape factors) — assignment minimum is 12
- **Target:** `Class` — 7 classes: BARBUNYA, BOMBAY, CALI, DERMASON, HOROZ, SEKER, SIRA
- **Missing values:** 0 · **Class counts after cleaning:** DERMASON 3546, SIRA 2636, SEKER 2027, HOROZ 1860, CALI 1630, BARBUNYA 1322, BOMBAY 522 (imbalance ratio 6.8:1, largest/smallest)

#	Feature	Description
1	`Area`	Number of pixels inside the bean region
2	`Perimeter`	Length of the bean border
3	`MajorAxisLength`	Longest line that can be drawn through the bean
4	`MinorAxisLength`	Longest line perpendicular to the major axis
5	`AspectRatio`	MajorAxisLength / MinorAxisLength (spelled `AspectRation` in the original Excel file)
6	`Eccentricity`	Eccentricity of the ellipse with the same moments as the region
7	`ConvexArea`	Pixels in the smallest convex polygon containing the bean
8	`EquivDiameter`	Diameter of a circle with the same area as the bean
9	`Extent`	Bounding-box pixels / bean area
10	`Solidity`	Convex-hull pixels / bean pixels (convexity)
11	`Roundness`	$4\pi A / P^2$ (spelled `roundness` in the original Excel file)
12	`Compactness`	EquivDiameter / MajorAxisLength
13	`ShapeFactor1`	Shape factor 1 (dimensionless)
14	`ShapeFactor2`	Shape factor 2
15	`ShapeFactor3`	Shape factor 3
16	`ShapeFactor4`	Shape factor 4

c. Dataset Source and Attribution

- UCI Machine Learning Repository: https://archive.ics.uci.edu/dataset/602/dry+bean+dataset (DOI https://doi.org/10.24432/C50S4B)
- Kaggle mirror: https://www.kaggle.com/datasets/muratkokludataset/dry-bean-dataset
- Authors' dataset page: https://www.muratkoklu.com/datasets/
- License: Creative Commons Attribution 4.0 (CC BY 4.0), so redistribution with attribution is permitted; the Excel file is bundled in `data/raw/`.
- **Citation:** KOKLU, M. and OZKAN, I.A., (2020), Multiclass Classification of Dry Beans Using Computer Vision and Machine Learning Techniques. Computers and Electronics in Agriculture, 174, 105507. DOI: https://doi.org/10.1016/j.compag.2020.105507
- UCI citation: *Dry Bean [Dataset]. (2020). UCI Machine Learning Repository. https://doi.org/10.24432/C50S4B*

*Two loaders are provided in the notebook (**Option A** UCI via `ucimlrepo`, **Option B** downloaded Kaggle/author file); both are converted to one standard schema by `src/data_loader.py`. Option B was used to produce all results in this repository.*

d. GitHub Repository Link

<YOUR_GITHUB_REPO_URL>

e. Data Preprocessing

1. Features and target separated; schema validated (16 features + `Class`).
2. Missing values: none found (the code would drop such rows; nothing is imputed or invented).
3. Duplicates: **68 exact duplicate rows removed** (first copy kept) *before* splitting, so no sample can appear in both train and test.
4. Target encoded with `LabelEncoder` (saved in `models/label_encoder.joblib`).
5. **Stratified 80:20 split**, `random_state=42` → **10,834 train / 2,709 test**, shared by all five models.
6. `StandardScaler` is the first step of every `Pipeline`, therefore fitted on training data only (no leakage; verified in `tests/run_checks.py`).
7. Outliers are kept (genuine measurements); nothing in the data is modified.

f. Models Used

Model	Pipeline	Key hyper-parameters (fixed a priori, not tuned on test data)
Logistic Regression	StandardScaler → LogisticRegression	`max_iter=2000`, L2, `C=1`
Decision Tree	StandardScaler → DecisionTreeClassifier	`max_depth=10`, `min_samples_leaf=5`
K-Nearest Neighbors	StandardScaler → KNeighborsClassifier	`n_neighbors=7`, `weights='distance'`
Naive Bayes	StandardScaler → GaussianNB	defaults
SVM	StandardScaler → SVC	RBF, `C=1`, `gamma='scale'`, `probability=True`

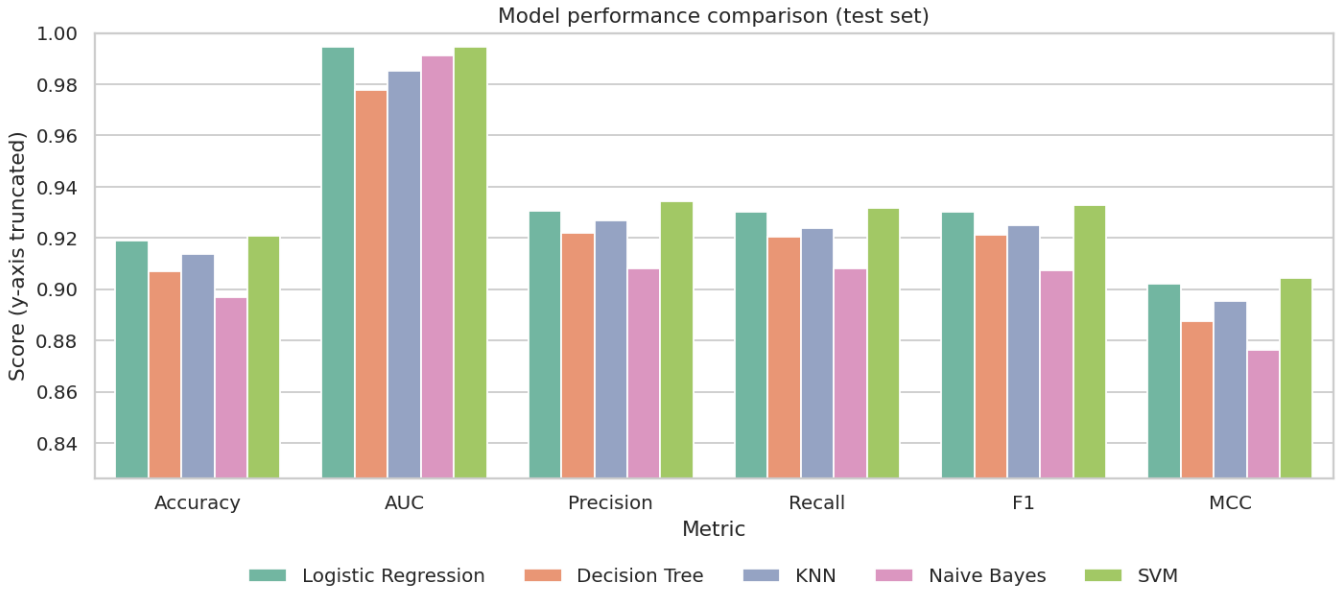
How each works, why it suits this data and its pros/cons are explained in the notebook (Phase 4).

g. Model Evaluation Metrics

Computed on the 2,709-row held-out test set; multiclass settings: **Accuracy**; **AUC** = one-vs-rest, macro-averaged, from `predict_proba`; **Precision / Recall / F1** = macro-averaged; **MCC** = multiclass Matthews correlation coefficient. `zero_division=0` handles classes that are never predicted.

h. Model Comparison Table

ML Model Name	Accuracy	AUC	Precision	Recall	F1	MCC
Logistic Regression	0.9192	0.9947	0.9307	0.9300	0.9302	0.9023
Decision Tree	0.9070	0.9779	0.9220	0.9204	0.9211	0.8875
KNN	0.9136	0.9854	0.9268	0.9237	0.9250	0.8956
Naive Bayes	0.8970	0.9911	0.9081	0.9082	0.9072	0.8762
SVM	0.9210	0.9945	0.9342	0.9317	0.9328	0.9045



Additional generalisation diagnostics (`reports/model_comparison_details.csv`; CV uses the training data only):

Model	Train acc.	5-fold CV acc.	Test acc.	Train–test gap
Logistic Regression	0.9260	0.9243 ± 0.0030	0.9192	0.0068
Decision Tree	0.9518	0.9112 ± 0.0059	0.9070	0.0448
KNN	1.0000	0.9240 ± 0.0055	0.9136	0.0864
Naive Bayes	0.8972	0.8968 ± 0.0060	0.8970	0.0002
SVM	0.9342	0.9300 ± 0.0057	0.9210	0.0132

i. Model Performance Observations

ML Model Name	Observation About Model Performance
Logistic Regression	Second-best on Accuracy (0.9192), Precision, Recall, F1 and MCC, and the highest AUC (0.9947), only 0.0002 above SVM. Train-test accuracy gap is just 0.0068, so it generalises very well. A linear decision surface on standardised features is almost as good as the non-linear SVM here.
Decision Tree	Accuracy 0.9070 and the lowest AUC of all models (0.9779): a tree outputs piecewise-constant leaf probabilities, so its class ranking is coarser. It shows the largest genuine overfitting among the models (train 0.9518 vs test 0.9070, gap 0.0448) even with max_depth=10 and min_samples_leaf=5.
KNN	Mid-table: Accuracy 0.9136, F1 0.9250, AUC 0.9854. Its training accuracy is 1.0000 only because distance weighting makes every training point its own nearest neighbour, so the 0.0864 train-test gap is partly an artefact; the 5-fold CV accuracy on the training data (0.9240) is the fairer indicator. Prediction is also the slowest at inference time because it stores the whole training set.
Naive Bayes	Weakest on Accuracy (0.8970), F1 (0.9072) and MCC (0.8762), yet third-highest on AUC (0.9911): it ranks classes reasonably but its hard decisions are worse. The independence assumption is clearly violated - 21 of 120 feature pairs have an absolute correlation above 0.9 - which is a likely reason. Gap 0.0002 means no overfitting; the limitation is bias, not variance.
SVM	Best on Accuracy (0.9210), Precision (0.9342), Recall (0.9317), F1 (0.9328) and MCC (0.9045); AUC (0.9945) is second, 0.0002 behind Logistic Regression. It also has the best 5-fold CV accuracy on the training data (0.9300) and a modest gap of 0.0132.

- **Best model overall by accuracy:** SVM · **Highest AUC:** Logistic Regression
- **Best precision/recall balance:** SVM (highest F1 = 0.9328; Precision 0.9342 and Recall 0.9317 are both the highest, differing by 0.0025)
- **Generalisation:** Naive Bayes has the smallest gap (0.0002) but at lower accuracy; among the strong models Logistic Regression (0.0068) and SVM (0.0132) generalise best.
- **Poor performers:** Naive Bayes (lowest Accuracy/F1/MCC); Decision Tree has the lowest AUC and the biggest real overfitting gap.
- Per-class (SVM): the hardest class is SIRA (F1 0.8684); BOMBAY is the easiest (F1 1.0000). The most frequent mix-up is DERMASON vs SIRA (96 test beans misassigned between them).

j. Overall Winning Model

Overall winner: Support Vector Machine (RBF). It ranks first on five of the six metrics (Accuracy, Precision, Recall, F1, MCC) and second on AUC (only 0.0002 behind Logistic Regression), and has the best cross-validated accuracy on the training data. Caveat: its accuracy lead over Logistic Regression is 0.0018, smaller than one standard error of the test accuracy (about 0.0052 for n=2709), so SVM and Logistic Regression should be regarded as statistically very close; Logistic Regression is the simpler, faster alternative.

k. Streamlit Application Description

`app.py` loads the five saved pipelines (it never retrains) and provides: a dashboard (dataset facts), CSV upload with validation (missing/unexpected columns, non-numeric or missing values, empty file, unknown class labels), a model dropdown, the six metrics, a confusion matrix (counts or row-normalised), a per-class classification report, a five-model comparison table + chart on the uploaded data, and a predictions table (actual vs predicted, confidence, misclassified filter) with CSV download. Each model has a short description in the sidebar, the confusion matrix shows both counts and row percentages, and the Evaluation tab states that uploaded data is not used for retraining. A bundled-sample button uses `data/test_data.csv` (the held-out test split). The uploaded file must contain the 16 feature columns and `Class`.

Saved-model / no-retraining deployment method. Training happens once, offline (`python -m src.train_models` or the notebook). Each model is saved as a complete scikit-learn `Pipeline` (StandardScaler + classifier) with `joblib.dump`, so the scaler statistics learned from the training split travel with the model. At run time `app.py` only calls `joblib.load`, `predict`, `predict_proba` and metric functions - it contains no `fit()` call (checked automatically in `tests/run_checks.py`) and does not import the training module. Uploaded rows are used only for prediction and scoring. Because pickled pipelines depend on library versions, `requirements.txt` pins `scikit-learn==1.8.0` and the other packages to the versions used for training.

l. Deployment Link

<YOUR_STREAMLIT_APP_URL>

m. Project Folder Structure

```
ML_Assignment_2/
├── README.md                   <- this file (also included in the submission PDF)
├── PROJECT_OVERVIEW.md         <- what every file does + how to run/deploy
├── SUBMISSION_CHECKLIST.md     <- the faculty guide's final checklist with honest status
├── requirements.txt            <- runtime dependencies (used by Streamlit Cloud)
├── requirements-dev.txt        <- adds jupyter/nbconvert for running the notebook
├── .gitignore
├── app.py                      <- Streamlit application
├── config.py                   <- paths, constants, dataset citation
├── .streamlit/config.toml      <- theme / upload limit
├── data/
│   ├── README.md
│   ├── raw/                    <- Dry_Bean_Dataset.xlsx (Option B) + dry_bean_raw.csv
│   ├── processed/              <- dry_bean_clean.csv, train.csv, test.csv
│   └── test_data.csv           <- the held-out 20 % test split (use it in the app)
├── notebooks/ML_Assignment_2.ipynb   <- executed notebook (Option A / B loaders, EDA, models)
├── src/
│   ├── data_loader.py          <- Option A (UCI) / Option B (file) loaders, validation
│   ├── preprocessing.py        <- cleaning, encoding, stratified split
│   ├── train_models.py         <- builds, trains, saves the five pipelines
│   ├── evaluate_models.py      <- six metrics, reports, figures
│   └── visualization.py        <- EDA + evaluation plots
├── models/                     <- 5 x .joblib pipelines + label_encoder.joblib + metadata.json
├── reports/
│   ├── model_comparison.csv, model_comparison_details.csv, TESTING_REPORT.md
│   ├── classification_reports/ <- per-model csv + txt
│   └── figures/                <- EDA + confusion matrices + comparison + ROC
├── screenshots/README.md       <- what to capture (BITS Virtual Lab etc.)
├── deployment/                 <- GITHUB_DEPLOYMENT.md, STREAMLIT_DEPLOYMENT.md
├── submission/                 <- ML_Assignment_2_Submission.pdf (template + README content)
├── tests/run_checks.py         <- automated test battery (writes reports/TESTING_REPORT.md)
└── tools/                      <- scripts that regenerate notebook / docs from the results
```

n. Installation Instructions

```
git clone <YOUR_GITHUB_REPO_URL>
cd ML_Assignment_2
python -m venv .venv
source .venv/bin/activate           # Windows: .venv\Scripts\activate
```

```
pip install -r requirements.txt      # add requirements-dev.txt if you want to run the notebook
```

Developed and tested with Python 3.12 (scikit-learn 1.8 requires Python ≥ 3.11). **scikit-learn is pinned to 1.8.0** because the `.joblib` files were saved with that version.

o. Execution Instructions

```
python -m src.train_models        # clean data, split, train the 5 pipelines, save to models/
python -m src.evaluate_models     # six metrics, reports, confusion matrices, figures
python -m src.visualization       # EDA figures
python tests/run_checks.py        # automated checks -> reports/TESTING_REPORT.md
streamlit run app.py              # launch the app, then upload data/test_data.csv
jupyter notebook notebooks/ML_Assignment_2.ipynb   # full walkthrough (needs requirements-dev.txt)
```

To use **Option A (UCI)** in the notebook, un-comment `data_loader.load_from_uci()` and comment out the Option B lines in the loading cell.

p. Conclusion

All five classifiers reach roughly 90%–92% test accuracy on the 7-class Dry Bean problem, with macro-AUC between 0.978 and 0.995. SVM (RBF kernel) gives the best overall result, with Logistic Regression statistically very close and much simpler; Naive Bayes is weakest, consistent with the strong feature correlations that violate its independence assumption. The hardest distinction is DERMASON vs SIRA. Limitations: hyper-parameters were not tuned, a single train/test split was used for the headline numbers, and differences between the top two models are within sampling noise. Possible extensions: grid/random search with nested cross-validation, feature selection to remove redundant size features, and ensemble methods.

Note: the results in this README are generated from the CSV files in `reports/`; `reports/TESTING_REPORT.md` lists what was and was not tested.